



Binary Classification of images in Python

This notebook shows the workflow for building, training, and testing a custom model class using Pytorch for Machine learning. The specific purpose is Binary Classification of images. This notebook covers up to the end of the training portion.

Training is implemented using Pytorch Ignite.

Project Overview:

The code in the cells of this notebook performs the following:

The images are first cleaned, duplicate checking is performed using the ImageHash package.

The data is visualized in various ways for inspection and inference.

The image files contained in the `train/test/val` folders have label files created which are stored in the main working directory.

The resulting file names will be:

```
["training_labels_final.csv", "val_labels_final.csv",  
"test_labels_final.csv"]
```

A transformer is initialized. Which performs various types of data augmentation.

```
## Transformer  
img_transforms = v2.Compose([v2.Resize((64,64)), ← Resizes the images.  
                             v2.RandomHorizontalFlip(p=0.5), ← Flips the images.  
                             v2.RandomPhotometricDistort(p=0.5), ← Distorts the images.  
                             v2.ToTensor())] ← Turns images into tensors.
```

Pytorch Custom Data set classes are then implemented. And have the following structure:

```
class Loading_training(Dataset):
    def __init__(self): ← Initializes the directory.
        """Loading various data sets
        but only in existing train, val test folders."""
    def __len__(self): ← Provides the number of samples in dataset.
    def __getitem__(self, idx): ← Loads samples from a given index.
```

A data loader is then used to 'feed' the files to the model.

```
train_dataloader = DataLoader(training, batch_size=batch_size, shuffle=True)
```

A custom Neural Network model is initialized.

```
class Modelo(nn.Module): ← The model class. A neural network.
    def __init__(self): ← Contains model layers.
        super(Modelo, self).__init__()
    def forward(self, x): ← Defines the order in which the layers are applied.
```

The basic model architecture is:

Layer (type)	Output Shape	Param #
-----	-----	-----
Conv2d-1	[-1, 64, 64, 64]	1,792
Conv2d-2	[-1, 64, 64, 64]	36,928
Conv2d-3	[-1, 64, 64, 64]	36,928
Conv2d-4	[-1, 64, 64, 64]	36,928
Conv2d-5	[-1, 32, 64, 64]	18,432
BatchNorm2d-6	[-1, 32, 64, 64]	0
MaxPool2d-7	[-1, 32, 32, 32]	0
LeakyReLU-8	[-1, 32, 32, 32]	0
Conv2d-9	[-1, 32, 32, 32]	9,248
Conv2d-10	[-1, 32, 32, 32]	9,248
Conv2d-11	[-1, 32, 32, 32]	9,248
Dropout2d-12	[-1, 32, 32, 32]	0
Conv2d-13	[-1, 16, 32, 32]	4,624
LeakyReLU-14	[-1, 16, 32, 32]	0
Conv2d-15	[-1, 16, 32, 32]	2,320
Conv2d-16	[-1, 16, 32, 32]	2,320

Conv2d-17	[-1, 16, 32, 32]	2,320
Conv2d-18	[-1, 8, 32, 32]	1,160
LeakyReLU-19	[-1, 8, 32, 32]	0
Dropout2d-20	[-1, 8, 32, 32]	0
Conv2d-21	[-1, 2, 32, 32]	146
MaxPool2d-22	[-1, 2, 16, 16]	0
LeakyReLU-23	[-1, 2, 16, 16]	0
Linear-24	[-1, 256]	131,328
Linear-25	[-1, 128]	32,896
Linear-26	[-1, 1]	129
=====		
Total params: 335,995		
Trainable params: 335,995		
Non-trainable params: 0		

Input size (MB): 0.05		
Forward/backward pass size (MB): 12.34		
Params size (MB): 1.28		
Estimated Total Size (MB): 13.67		

A loss function and optimizer are chosen.

```
loss_fn_ = nn.BCEWithLogitsLoss() ← Combines sigmoid layer with BCE Loss.
optimizer = torch.optim.AdamW(model.parameters(), lr = lr) ← Specifies learning rate
(and other params).
```



Pytorch Ignite is used for model training and to create model checkpoints.

```
def update_model(engine, batch): ← Creates a basic trainer.
    return x()
trainer = Engine(update_model) ← Loops over the training data.

val_metrics = { ← Metrics to be tracked.
    "metric": metric(),
}

def validation_step(engine, batch): ← Runs model validation.
    return x()
evaluator = Engine(validation_step) ← Loops over the validation data.

@trainer.on(Events.ITER_COMPLETED(every=x)) ← How often to log events(iteration).
def log_training(engine):
```

```

@trainer.on(Events.EPOCH_COMPLETED(every=x)) ← How often to log events(per X epochs).
def run_validation():
    evaluator.run(val_dataloader)

def score_function(engine): ← Returns metrics that were defined in val_metrics.
    return engine.state.metrics

model_checkpoint = ModelCheckpoint("checkpoint") ← Creates a checkpoint.

trainer.run(train_dataloader, max_epochs=X) ← Runs the trainer.

```

The checkpoint models are then tested to see which performs best against unseen data from the original dataset as well as data from a completely new dataset.

```

for j in os.listdir(X) ← Iterates over all of the models.
    PATH=os.path.join(X,j)
    model.load_state_dict(torch.load(PATH)) ← Loads state dictionaries.
    for i,batch in enumerate(test_dataloader): ←Performs testing.
        . . .

```

Code:

Installing libraries and dependencies.

Python version 11.3.7

```

#Installing libraries and dependencies

import os
import pandas as pd
import numpy as np
import torch
from ignite.handlers import Checkpoint
from PIL import Image
from torchvision.transforms import v2

```

```

from torch.utils.data import Dataset
from torchvision.transforms import ToTensor
from torchvision.io import read_image
from torch import nn
from torch.utils.data import DataLoader
from torchvision.transforms import Compose, Normalize, ToTensor
from ignite.engine import Engine, Events, create_supervised_trainer,
create_supervised_evaluator
from ignite.metrics import Accuracy, Loss
from ignite.handlers import ModelCheckpoint
from ignite.contrib.handlers import TensorboardLogger, global_step_from_engine
from ignite.handlers import Timer, BasicTimeProfiler, HandlersTimeProfiler
import warnings
warnings.filterwarnings('ignore')

```


Visualizations and clean up

```

larger_list=[]
directory = os.path.join(os.getcwd(), 'train')
q = 1
fig,ax = plt.subplots(figsize=(15, 15))
for I in range(0,(25)):
    random_file = random.choice(os.listdir(directory))
    abc = (os.path.join(directory, random_file)) #<- get the first file name
    picture = Image.open(abc) #<- reading those images
    plt.subplot(5,5,q) #<- you know that is so neat I am leaving it in
    plt.axis('off')
    title = ("Who doesn't like binary classification of cute animals")
    plt.imshow(picture)
    plt.suptitle(title)
    ax.set_axis_off()
    q+= 1 #<- Make it so there are not so many files
    #print(q) #<- So I know it is not running forever
    if q == 25:
        break

```

Who doesn't like binary classification of cute animals



Checking for duplicate values using hashing

The `average_hash` function was used from the Imagehash package.

<https://pypi.org/project/ImageHash/>

The image is first scaled down, then changed to grayscale, the means of the image colors are computed. The bits are then computed. A hash is then constructed by turning the 64 bits into a 64 bit integer. And if the results are the same, the images are likely the same.

<https://www.hackerfactor.com/blog/index.php/?archives/432-Looks-Like-It.html>

The code below makes lists of any possible duplicates in the files. And then erases them.

```
list_for_hashing = []
files_to_keep = []
files_to_purge = []
files_to_purge1 = []
for file in os.listdir(directory1):
    image1=imagehash.average_hash(Image.open(os.path.join(directory1,file)))
    if image1 not in list_for_hashing:
        list_for_hashing.append(image1)
        files_to_keep.append(file)
```

```

else:
    print(file,image1)
    files_to_purge.append([file,image1])
    files_to_purge1.append(image1)

```

```

zipped = list(zip(files_to_keep, list_for_hashing))
m=pd.DataFrame(zipped)
q = pd.DataFrame(files_to_purge)
pqm= q.merge(m, on = 1, how = 'inner')
file_list = list(pqm['0_x'])
file_list1 = list(pqm['0_y'])

```

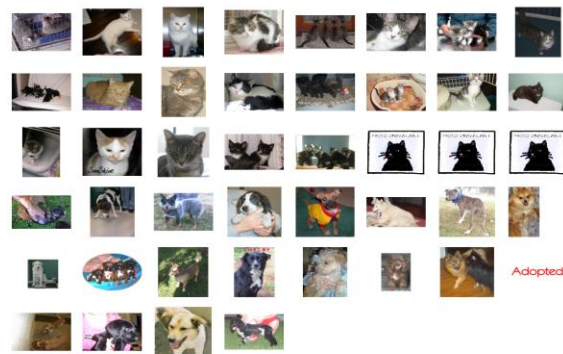
This code plots all of the duplicate images

```

from matplotlib import pyplot as plt
larger_list=[]
directory= (os.getcwd()+'\PetImages\Cat')
q = 1
lengths1=len(pqm)
fig,ax = plt.subplots(figsize=(15,15))
for i in os.listdir(directory):
    if i in file_list:
        abc=(os.path.join(directory, i))
        picture = Image.open(abc) #<- reading those images
        plt.subplot(8,8,q) #<- you know that is so neat I am leaving it in
        plt.axis('off')
        title =("Who doesn't like binary classification of cute animals")
        plt.imshow(picture)
        plt.suptitle(title)
        ax.set_axis_off()
        q+=1 #<- Make it so there are not so many files

```

Who doesn't like binary classification of cute animals



This portion of code erases the duplicates

```
directory= (os.getcwd()+'\train')
for file in os.listdir(directory):
    if file in file_list1:
        okay= os.path.join(directory,file)
        if os.path.exists(okay):
            os.remove(okay)
    else:
        pass
```

```
from matplotlib import pyplot as plt
larger_list=[]
directory= (os.getcwd()+'\train')
q = 1
lengths1=len(pqm)
fig,ax = plt.subplots(figsize=(15,15))
for i in os.listdir(directory):
    if i in file_list:
        abc=(os.path.join(directory, i))
        picture = Image.open(abc) #<- reading those images
        plt.subplot(5,5,q) #<- you know that is so neat I am leaving it in
        plt.axis('off')
        title =("Who doesn't like binary classification of cute animals")
```



```
plt.imshow(picture)
plt.suptitle(title)
ax.set_axis_off()
q+=1 #<- Make it so there are not so many files
```

The result is a blank plot. So I will not show it here.

Making inferences about the image data, such as pixel density, the sizes of the images and the overall balance of the dataset.

```
pixel_values1 = []
widths_cat=[]
heights_cat=[]
widths_dog=[]
heights_dog=[]
traindir= directory
partial_string = 'og'
for img in os.listdir(traindir):
    if 'og' in img:
        img_path = os.path.join(traindir +'\\'+img ) # Making image file path
        im = Image.open(img_path)
        widths_dog.append(im.size[0])
        heights_dog.append(im.size[1])
        pixel_values1 = list(im.getdata())
    else:
        img_path = os.path.join(traindir +'\\'+img ) # Making image file path
        im = Image.open(img_path)
        widths_cat.append(im.size[0])
        heights_cat.append(im.size[1])
        pixel_values2 = list(im.getdata())
```

```

pixel_vals_b_dog = [x[0] for x in pixel_values1]
pixel_vals_g_dog = [x[1] for x in pixel_values1]
pixel_vals_r_dog = [x[2] for x in pixel_values1]
pixel_vals_b_cat = [x[0] for x in pixel_values2]
pixel_vals_g_cat = [x[1] for x in pixel_values2]
pixel_vals_r_cat = [x[2] for x in pixel_values2]

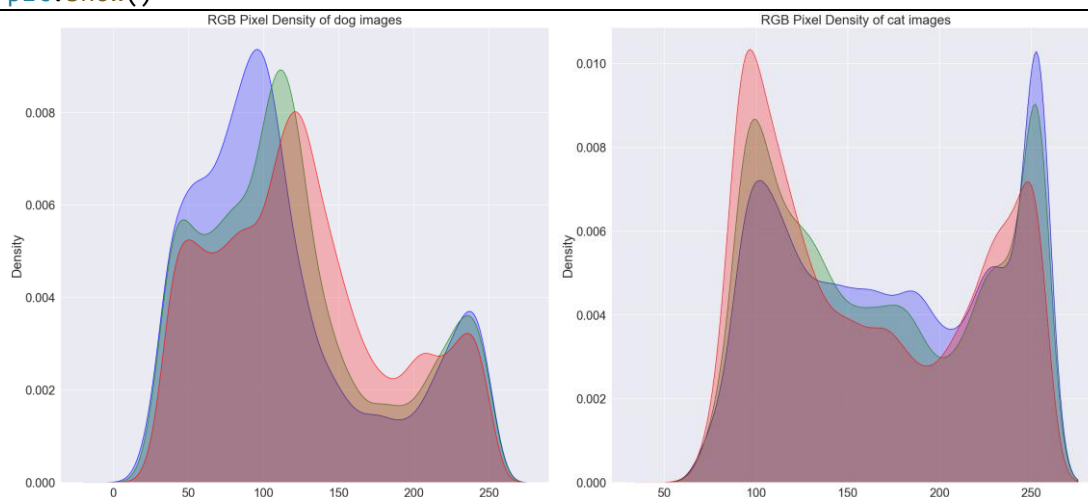
```

```

f,ax= plt.subplots(figsize=(16,15))
sns.kdeplot(pixel_vals_b_dog, ax=ax,color='blue',fill=True).set(title='RGB Pixel
Density of dog images')
sns.kdeplot(pixel_vals_g_dog, ax=ax,color='green', fill = True)
sns.kdeplot(pixel_vals_r_dog, ax=ax,color='red', fill =True)

f,ax= plt.subplots(figsize=(16,15))
sns.kdeplot(pixel_vals_b_cat, ax=ax,color='blue',fill=True).set(title='RGB Pixel
Density of cat images')
sns.kdeplot(pixel_vals_g_cat, ax=ax,color='green', fill = True)
sns.kdeplot(pixel_vals_r_cat, ax=ax,color='red', fill =True)
plt.show()

```



Images are composed of pixels. Pixels store information about their colors, in the form of 3 values. Red, green, blue. (R, G, B) When a package like `imread` is used to process an image, the resulting data is stored in a tuple (list?) where the first value is R, the second is G,

the third is B. And if the values from each section are taken and plotted it can provide insight as to the general composition of the images in each category.

```
# Pixel values need to be in range 0/1
pixel_values1 = [50/255, 50/255, 50/255]
pixel_values2 = [100/255, 100/255, 100/255]
pixel_values3 = [150/255, 150/255, 150/255]
def plot_pixels(pixel_values):
    patch = plt.Rectangle((0, 0), 1, 1, color = pixel_values)
    fig, ax = plt.subplots()
    ax.add_patch(patch)
    ax.set_xlim(0, 1)
    ax.set_ylim(0, 1)
    ax.axis('off')
plot_pixels(pixel_values1), plot_pixels(pixel_values2), plot_pixels(pixel_values3)
plt.show()
```



There is a lot of density in the 50, 100, and 150 range for both sets of images. And when the colors are plotted, voilà. Many of the images contain large proportions of neutral colors.

What about size distributions between the images? Are dog or cat images generally larger?

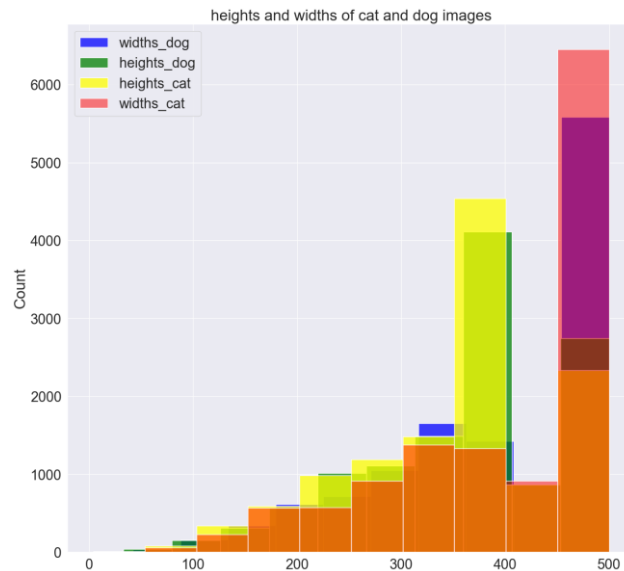
```
f,ax= plt.subplots(figsize=(16,15))
sns.kdeplot(widths_dog, ax=ax,color='blue',fill=True).set(title=
'heights and widths of cat and dog images')
sns.kdeplot(heights_dog, ax=ax,color='green', fill = True)
sns.kdeplot(heights_cat, ax=ax,color='yellow', fill = True)
```

```
sns.kdeplot(widths_cat, ax=ax,color='red', fill = True)
plt.legend(['widths_dog','heights_dog','heights_cat','widths_cat'],loc='upper left')
plt.show()
```



Most images for both categories are generally on the larger side. Most heights cluster around 350 pixels for both cat and dog images. While many of the images are about 500 pixels high.

```
f,ax= plt.subplots(figsize=(16,15))
sns.histplot(widths_dog, ax=ax,color='blue',fill=True, bins =10).set(title=
'heights and widths of cat and dog images')
sns.histplot(heights_dog, ax=ax,color='green', fill = True,bins=10)
sns.histplot(heights_cat, ax=ax,color='yellow', fill = True,bins=10)
sns.histplot(widths_cat, ax=ax,color='red', fill = True,alpha= 0.5,bins=10)
plt.legend(['widths_dog','heights_dog','heights_cat','widths_cat'],loc='upper left')
plt.show()
```



The same data visualized as a histogram shows much the same pattern. Almost half of the 24,000 images are 500 pixels wide. A set of subplots helps illustrate this better.

```
f,axes = plt.subplots(1, 4,figsize=(20,10))
cats_heights1 = pd.DataFrame(heights_cat,columns=['heights_cat'])
cats_width1 = pd.DataFrame(widths_cat,columns=['widths_cat'])
dogs_heights1=pd.DataFrame(heights_dog,columns=['heights_dog'])
dogs_width1=pd.DataFrame(widths_dog,columns=['widths_dog'])
sns.histplot(cats_heights1, x="heights_cat",ax=axes[0],
color='red').set(title="heights of cat images")
sns.histplot(cats_width1, x="widths_cat",ax=axes[1],color='blue').set(title="widths
of cat images")
sns.histplot(dogs_heights1,
x="heights_dog",ax=axes[2],color='green').set(title="heights of dog images")
sns.histplot(dogs_width1, x="widths_dog",ax=axes[3]).set(title="widths of dog
images")
```



The subplots help illustrate the spread of image sizes. All sets of images have one category that is as large as the sum of all of the others. The largest categories for heights are the same for both cat and dog images. About 400 pixels and 500 respectively. And both have similar counts for each, about 3500 images at about 400 pixels high and about 2000 at 500 pixels for cats. Versus about 3000 images at 400 pixels and about 2000 images at 500 pixels for dogs.

Widths exhibit a similar pattern, both cat and dog images cluster at about 500 pixels for the most frequent category. With cat images at this size numbering almost 6000, and dog images trailing closely behind at about 5000.

What about the balance of samples? How many classes are found in each data set?

Since the problem is binary, the images fall into categories of 0 or 1.

```
print(' There are:', len(xy1['status'].unique()), 'categories in the training data set', '\n',
      'There are:', len(xy2['status'].unique()), 'categories in the validation data set', '\n',
      'There are:', len(xy3['status'].unique()), 'categories in the test data set')
```

There are: 2 categories in the training data set
There are: 2 categories in the validation data set
There are: 2 categories in the test data set

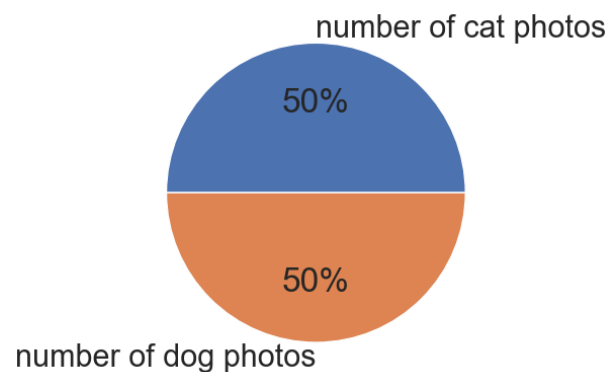
Not surprisingly, there are two categories in each data set.
How about the numbers of samples?

```
print('Training data:', 'images in the 0 category : ',len(xy1[xy1['status']==0]),",",  
      , 'images in the 1 category : ',len(xy1[xy1['status']==1]))  
print('Validation data:', 'images in the 0 category:',len(xy2[xy2['status']==0]),",",  
      , 'images in the 1 category:',len(xy2[xy2['status']==1]))  
print('Test data:', 'images in the 0 category:',len(xy3[xy3['status']==0]),",",  
      , 'images in the 1 category:',len(xy3[xy3['status']==1]))
```

Training data: images in the 0 category : 12418 , images in the 1 category : 12418
Validation data: images in the 0 category: 40 , images in the 1 category: 40
Test data: images in the 0 category: 40 , images in the 1 category: 40

To help drive home that the data set contains equal numbers. A pie chart is included.

```
data=[len(heights_cat), len(heights_dog)]  
keys='number of cat photos','number of dog photos'  
palette_color = sns.color_palette('light')  
plt.pie(data, labels=keys, colors = palette_color,autopct='%.0f%%')  
plt.show()
```



Unsurprisingly, the pie is split evenly.

Procedure

Setting the device to cuda, so the GPU can be utilized.

```
print(torch.__version__)  
device = torch.device('cuda')
```

Creating lists, one for training, validation and testing, appending relevant file names to each.

```
# Loading data  
# For training  
  
starter_path = os.path.join(os.getcwd(), 'train')  
data_labels_tr = []  
next_path = os.listdir(starter_path)  
  
for i in next_path:  
    impath = (os.path.join(starter_path, i))  
    data_labels_tr.append(i)  
  
#For validation  
starter_path = os.path.join(os.getcwd(), 'val')  
data_labels_val = []  
next_path = os.listdir(starter_path)  
  
for i in next_path:  
    impath = (os.path.join(starter_path, i))  
    data_labels_val.append(i)  
  
#For testing  
starter_path = os.path.join(os.getcwd(), "test")  
data_labels_tst = []  
next_path = os.listdir(starter_path)  
  
for i in next_path:
```



```
impath = (os.path.join(starter_path, i))
data_labels_tst.append(i)
```

Creating labels by matching image names to partial strings. If the image contains dog or in this case 'og' a 1 will be added to a new column titled 'status'. Column names are also added to simplify my life.

```
# A lazyish way of creating labels for a binary classification task. Put a 0 if the #
# file contains a partial name
partial_string = 'og'
xy1=pd.DataFrame(data_labels_tr)
xy2=pd.DataFrame(data_labels_val)
xy3=pd.DataFrame(data_labels_tst)
xy1['status'] = xy1[0].str.contains(partial_string).astype(int)
xy2['status'] = xy2[0].str.contains(partial_string).astype(int)
xy3['status'] = xy3[0].str.contains(partial_string).astype(int)
xy1.columns=['id', 'status']
xy2.columns=['id', 'status']
xy3.columns=['id', 'status']
```

All data is saved to csv files, which **will** overwrite whatever is in the directory if it exists.

```
# Save to csv files in the current directory
xy1.to_csv("training_labels_final.csv",index=False)
xy2.to_csv("val_labels_final.csv",index=False)
xy3.to_csv("test_labels_final.csv",index=False)
```

Resizing images, flipping them and turning them into tensors.

```
## Transformer
img_transforms = v2.Compose([v2.Resize((64,64)),
                             v2.RandomHorizontalFlip(p=0.5),
                             v2.RandomPhotometricDistort(p=0.5),
                             v2.ToTensor()] )
```

Custom Data loaders load all labels, pair them to images and apply transformations.

```
# Custom data loaders for training validation testing
# Loads all images and puts them into either lists or into data dictionaries.

class Loading_training(Dataset):

    def __init__(self):
        """Loading various data sets
        but only in existing train, val test folders."""
        self.selected_dataset_dir = os.path.join(os.path.join(os.getcwd(), "train"))
        self.all_filenames = os.listdir(self.selected_dataset_dir)
        self.all_labels = pd.read_csv(os.path.join(os.getcwd(),
'training_labels_final.csv'))
        self.label_meanings = self.all_labels.columns.values.tolist()

    def __len__(self):
        """Weird"""
        return len(self.all_filenames)

    def __getitem__(self, idx):
        img_path = os.path.join(self.selected_dataset_dir, self.all_labels.iloc[idx,
0])

        image = Image.open(img_path).convert("RGB")
        label = self.all_labels.iloc[idx, 1]
        label = torch.tensor(label)
        image = img_transforms(image)
        return image, label

class Loading_val(Dataset):

    def __init__(self):
        """Trying to load jpeg transformed files."""
```

```

        self.selected_dataset_dir = os.path.join(os.getcwd(), "val")
        self.all_filenames = os.listdir(self.selected_dataset_dir)
        self.all_labels = pd.read_csv(os.path.join(os.getcwd(),
'val_labels_final.csv'))
        self.label_meanings = self.all_labels.columns.values.tolist()

    def __len__(self):
        """Weird"""
        return len(self.all_filenames)

    def __getitem__(self, idx):
        img_path = os.path.join(self.selected_dataset_dir, self.all_labels.iloc[idx,
0])

        image = Image.open(img_path).convert("RGB")
        label = self.all_labels.iloc[idx, 1]
        label = torch.tensor(label)
        image = img_transforms(image)
        return image, label

class Loading_test(Dataset):

    def __init__(self):
        """Loading test images."""
        self.selected_dataset_dir = os.path.join(os.getcwd(), "test")
        self.all_filenames = os.listdir(self.selected_dataset_dir)
        self.all_labels =
pd.read_csv(os.path.join(os.getcwd(), 'test_labels_final.csv'))
        self.label_meanings = self.all_labels.columns.values.tolist()

    def __len__(self):
        """Weird"""
        return len(self.all_filenames)

```


```

def __getitem__(self, idx):
    """The IDX function is built in to Pytorch. Pretty Cool I think.
    And then I use the above image transforms function"""
    selected_filename = self.all_filenames[idx]
    #print(selected_filename)
    imagepil = Image.open(os.path.join(self.selected_dataset_dir,
selected_filename)).convert("RGB")
    image = img_transforms(imagepil)
    self.all_labels = pd.read_csv(os.path.join(os.getcwd(),
'test_labels_final.csv'))
    label= self.all_labels['status'][idx]

    sample = {'data':image,
              'label':label,
              'img_idx':idx, 'sample_name':selected_filename}
    return(sample)

```

The model accepts 64 X 64 3 channel RGB image as input, and casts a prediction about the class of the image as output.



Several convolutional layers are stacked together at a time into blocks. The consistent padding throughout the model maintains sizes. Layer size can be calculated using the following formula:

$$\frac{I - K + (2 * P)}{S} + 1$$

- Where I is the shape of your input.
- K is the filter size. (The kernel size)
- P is the padding of the input.
- S is the stride.

The first input size is 64 X 64 X 3, and this shape is maintained for almost the entire block.

A calculation for the first layer shape as an example:

$$\frac{64 - 3 + (2 * 1)}{1} + 1 = 64$$

The second block scales outputs to 32 X 64 X 3 then further down to 32 X 32 X 3. The third block, to 16 X 16 X 3 and the final block uses dimensions of 8 X 8 X 3. With a single final layer to scale the output down to an appropriate number of channels (2).

Within the model layers MaxPooling is also used for scaling down inputs. Each application reduces the size by a factor of 2.

First MaxPool = 64 / 2 = 32

Second MaxPool = 32 / 2 = 16

LeakyRelu is used throughout the model architecture. It is thought that LeakyRelu helps to combat the vanishing gradient problem. The final layers of the model are linear.

The first input is (2 X 16 X 16) or 512 features. Which is then divided by 2 and then further subdivided by 2. The final layer outputs the prediction.

First Linear Layer → 512 features → 256 features

Second Linear Layer → 256 features → 128 features

Third Linear Layer → 128 features → 1 feature

The loss function is Binary Cross Entropy with Logits Loss.
The Optimizer is AdamW.

Model Architecture:

```
class Modelo(nn.Module):

    def __init__(self):
        super(Modelo, self).__init__()
        self.main = torch.nn.Sequential(
            torch.nn.Conv2d(in_channels=3, out_channels=64, kernel_size=(3, 3),
padding=1),
            torch.nn.Conv2d(in_channels=64, out_channels=64, kernel_size=(3, 3),
padding=1),
            torch.nn.Conv2d(in_channels=64, out_channels=64, kernel_size=(3, 3),
padding=1),
            torch.nn.Conv2d(in_channels=64, out_channels=64, kernel_size=(3, 3),
padding=1),
            #torch.nn.Dropout2d(p=0.2),
            torch.nn.Conv2d(in_channels=64, out_channels=32, kernel_size=(3, 3),
padding=1,bias=False),
            nn.BatchNorm2d(32, affine=False),
            torch.nn.MaxPool2d(2, 2),
            torch.nn.LeakyReLU(inplace=True,negative_slope=0.02))
        self.main1 = torch.nn.Sequential(
            torch.nn.Conv2d(in_channels=32, out_channels=32, kernel_size=(3, 3),
padding=1),
            torch.nn.Conv2d(in_channels=32, out_channels=32, kernel_size=(3, 3),
padding=1),
            torch.nn.Conv2d(in_channels=32, out_channels=32, kernel_size=(3, 3),
padding=1),
            torch.nn.Dropout2d(p=0.2),
            torch.nn.Conv2d(in_channels=32, out_channels=16, kernel_size=(3, 3),
padding=1),
            torch.nn.LeakyReLU(inplace=True,negative_slope=0.02))
        self.main2 = torch.nn.Sequential(
```

```

        torch.nn.Conv2d(in_channels=16, out_channels=16, kernel_size=(3, 3),
padding=1),
        torch.nn.Conv2d(in_channels=16, out_channels=16, kernel_size=(3, 3),
padding=1),
        torch.nn.Conv2d(in_channels=16, out_channels=16, kernel_size=(3, 3),
padding=1),
        torch.nn.Conv2d(in_channels=16, out_channels=8, kernel_size=(3, 3),
padding=1),
        #torch.nn.MaxPool2d(2, 2),
        torch.nn.LeakyReLU(inplace=True,negative_slope=0.02))
    self.main3 = torch.nn.Sequential(
        torch.nn.Dropout2d(p=0.2),
        torch.nn.Conv2d(in_channels=8, out_channels=2, kernel_size=(3, 3),
padding=1),
        torch.nn.MaxPool2d(2, 2),
        torch.nn.LeakyReLU(inplace=True,negative_slope=0.1))
    #self.fc1 = nn.Linear(in_features=8192,out_features=4096)
    #self.fc2 = nn.Linear(in_features=4096,out_features=2048)
    #self.fc3 = nn.Linear(in_features=2048,out_features=1024)
    self.fc4 = nn.Linear(in_features=512,out_features=256)
    self.fc5 = nn.Linear(in_features=256,out_features=128)
    self.fc6 = nn.Linear(in_features=128,out_features=1)

def forward(self, x):
    x = self.main(x)
    x = self.main1(x)
    x = self.main2(x)
    x = self.main3(x)
    x = torch.flatten(x,start_dim=1,end_dim=-1)
    #x = self.fc1(x)
    #x = self.fc2(x)
    #x = self.fc3(x)
    x = self.fc4(x)
    x = self.fc5(x)
    x = self.fc6(x)
    return x

```

```
model = Modelo().to(device)

training1=Loading_training()
val1=Loading_val()
test1= Loading_test()

train_dataloader = DataLoader(training1, batch_size=64, shuffle=True)
val_dataloader = DataLoader(val1, batch_size=10, shuffle=True)
```

```
loss_fn=nn.BCEWithLogitsLoss()
optimizer = torch.optim.AdamW(model.parameters(), lr=0.0000050)
```

The trainer below is adapted from Pytorch Ignite.

```
def update_model(engine, batch):
    model.train()
    data,label = batch
    optimizer.zero_grad()
    data=data.to(device)
    label=label.to(device)
    #print(label.get_device())
    #print(data.get_device())
    outputs,_=(model(data),label)
    outputs=outputs.squeeze()
    loss = loss_fn(outputs, label.float())
    loss.backward()
    optimizer.step()
    return loss.item()

trainer = Engine(update_model)
```



```
val_metrics = {
    "accuracy": Accuracy(),
    "loss": Loss(loss_fn_)
}

def validation_step(engine, batch):
    model.eval()
    with torch.no_grad():
        x, y = batch
        x=x.to(device)
        y=y.to(device)
        outputs,_=(model(x),y)
        outputs = outputs.squeeze()
        outputs=torch.sigmoid(outputs)
        outputs=outputs.round()
        y=y.cpu().detach()

    return outputs, y.float()

evaluator = Engine(validation_step)

from ignite.metrics import Accuracy
from ignite.metrics import Precision, Recall

precision = Precision()

Accuracy().attach(evaluator, "accuracy")
Precision().attach(evaluator, 'precision')
Recall(average='weighted').attach(evaluator, 'recall')
```

```

@trainer.on(Events.ITERATION_COMPLETED(every=100))
def log_training(engine):
    batch_loss = engine.state.output
    lr = optimizer.param_groups[0]['lr']
    e = engine.state.epoch
    n = engine.state.max_epochs
    i = engine.state.iteration
    print("Epoch {}/{} : {} - batch loss: {}, lr: {}".format(e, n, i, batch_loss,
lr))

from ignite.handlers import Timer, BasicTimeProfiler, HandlersTimeProfiler
from ignite.engine import Events

validate_every = 5

@trainer.on(Events.EPOCH_COMPLETED(every=validate_every))
def run_validation():
    evaluator.run(val_dataloader)

@trainer.on(Events.EPOCH_COMPLETED(every=validate_every))
def log_validation():
    ugh=[]
    metrics = evaluator.state.metrics
    print(f"Epoch: {trainer.state.epoch}, Accuracy:
{metrics['accuracy']}, Precision: {metrics['precision']}, recall:
{metrics['recall']}")

def score_function(engine):
    return engine.state.metrics["accuracy"]

model_checkpoint = ModelCheckpoint(
    "checkpoint",

```

```

        n_saved=25,
        filename_prefix="best",
        score_function=score_function,
        score_name="accuracy",
        global_step_transform=global_step_from_engine(trainer),
    )

evaluator.add_event_handler(Events.COMPLETED, model_checkpoint, {"model": model})

tb_logger = TensorboardLogger(log_dir="tb-logger")

tb_logger.attach_output_handler(
    trainer,
    event_name=Events.ITERATION_COMPLETED(every=100),
    tag="training",
    output_transform=lambda loss: {"batch_loss": loss},
)

for tag, evaluator in [("training", trainer), ("validation", evaluator)]:
    tb_logger.attach_output_handler(
        evaluator,
        event_name=Events.EPOCH_COMPLETED,
        tag=tag,
        metric_names="all",
        global_step_transform=global_step_from_engine(trainer),
    )

trainer.run(train_dataloader, max_epochs=200)

```

Model iterating:

```

Epoch 1/200 : 100 - batch loss: 0.6717314720153809, lr: 5e-06
Epoch 1/200 : 200 - batch loss: 0.62602299451828, lr: 5e-06
Epoch 1/200 : 300 - batch loss: 0.5983356833457947, lr: 5e-06
Epoch 2/200 : 400 - batch loss: 0.620399534702301, lr: 5e-06
Epoch 2/200 : 500 - batch loss: 0.6466283798217773, lr: 5e-06
Epoch 2/200 : 600 - batch loss: 0.6703633069992065, lr: 5e-06
Epoch 2/200 : 700 - batch loss: 0.6332987546920776, lr: 5e-06
Epoch 3/200 : 800 - batch loss: 0.5965307950973511, lr: 5e-06
Epoch 3/200 : 900 - batch loss: 0.5554847121238708, lr: 5

```